

ClipQuest answer grading & question quality

Release QA report • 18 August 2026 • build 6674041e • Worker 8072a257

Outcome

The repaired validator completed 8 of 10 cached YouTube sources, generating 80 questions. Every accepted learner answer graded correct, every negative-control answer was rejected, and all 26 short-answer canonical responses passed. Two sources were deliberately blocked by the quality gates instead of allowing a repeated objective or an incomplete placeholder answer into a quiz.

Metric	Observed	Interpretation
Sources tested	10 (8 pass / 2 blocked)	Blocked output is safer than shipping malformed content.
Questions accepted	80	MC, true/false, and short answer were exercised.
Positive controls	80	Canonical answers graded correct.
Negative controls	80	Wrong MC/TF and "I do not know" short answers rejected.
Short answers	26 / 26 correct	Accepted-answer alternatives were also checked where present.
Local DeepSeek calls	85	Generation stayed local; no captions were sent to ClipQuest Worker AI.
Automatic retries	5	Bounded recovery ran without weakening quality validation.

What changed

Reason-first local grading

Added a shared LocalAnswerGrade contract and a DeepSeek function tool. The model receives only the learner-facing question, the human response, type, and options; it writes a short reason first, then calls grade_answer with correct/wrong, confidence, and matched ideas.

Softer semantic feedback

Natural paraphrases, concise fragments, and meaning-equivalent short answers are accepted by the local grader. The API score remains authoritative for persistence so a modified browser cannot rewrite mastery.

True/false polarity

Kept locally assigned polarity authoritative and rejected false items that simply repeat the supported true claim. This prevents a polarity label from contradicting the statement.

Exact-stem protection

Repeated concepts remain allowed, but prompt-first generation now blocks an exactly repeated learner-facing stem against accepted questions. Near-identical grading targets are still retried.

Answer-to-prompt consistency

Added checks for condition prompts that lack a condition, why/how prompts that only name a term, and short-answer targets that are not complete responses.

Fragment and absolute wording gates

Rejects dangling conjunctions and placeholders such as “another effect” or “a third consequence”; rejects unsupported “always/never/all/none/every/must” claims while allowing source-supported absolutes.

Cross-platform bridge

Added answer-grading-v1 to the Chrome extension, web bridge, iOS, and Android local clients. Extension and app bundle rebuilt as 0.8.20.

Ten-source grading matrix

The matrix below is generated from the post-fix local DeepSeek run. It records counts and validation outcomes only; it intentionally omits transcript text and API credentials.

Video ID	Status	Questions	Types (MC / TF / short)	Positives	Negatives rejected	Short answers	Retries
dJplU1rSOFY	passed	15	5/5/5	15	15	5/5	5
rcOFV4y5z8c	passed	10	3/4/3	10	10	3/3	0
q5qaGHfdmYM	passed	15	5/5/5	15	15	5/5	0
CGoNpwN0mrs	passed	5	2/2/1	5	5	1/1	0
eQEaiZ2j9oc	failed	0	0/0/0	0	0	0/0	0
nBO7zWO1zXY	passed	15	5/5/5	15	15	5/5	0
HQ-Kg_xgdhE	passed	5	2/1/2	5	5	2/2	0
yubzJw0uiE4	passed	10	4/3/3	10	10	3/3	0
Sfzo4xm5eX8	failed	0	0/0/0	0	0	0/0	0
Og5xAdC8EUI	passed	5	1/2/2	5	5	2/2	0

Blocked-source diagnostics

eQEaiZ2j9oc: The grading target repeats an already accepted objective.

Sfzo4xm5eX8: The learner-facing question and answer are not complete and well-supported.

These are quality failures, not accepted quiz defects. The correct behavior is to retry with a different grounded objective or show recovery/configuration guidance; shipping the malformed candidate would violate the requested “not exactly the same,” prompt/answer consistency, and incomplete-answer rules.

Live Chrome and Cloudflare evidence

- Chrome production tab: <https://clipquest.ccwu.cc/> (signed-in session).
- Live DOM showed completed cards with enabled “Export notes” and incomplete cards with disabled “Notes not ready.”
- Clicked the enabled export action for “The History of Femboys”; a PDF was downloaded as “The-History-of-Femboys-cheat-sheet (2).pdf”.
- Cloudflare Worker version 8072a257-adaf-4706-afd7-fce6587c4bdc was promoted to 100% traffic with tag 6674041e4959185b51311d13dfe3810391a01d83.
- Production asset probe at +0s passed: 9 HTML shells and 9 entry bundles, with version affinity present.
- Previous release long-tail probes also passed at +0s, +120s, +300s, and +600s for the first deployment.

Chrome extension note

The existing Chrome profile had an older loaded extension during the live tab check, so the new answer-grading-v1 capability was verified in source/tests and in the published 0.8.20 ZIP, but not claimed as loaded into that profile. Installing/reloading a browser extension is a separate browser-side action requiring an action-time confirmation; the live export path itself was verified successfully.

Remaining issues and next actions

- Two of ten cached sources remain blocked after bounded retries: one repeats an accepted objective and one emits an incomplete placeholder answer. They are not silently accepted.
- The local DeepSeek reason is learner-facing feedback; persisted correctness/mastery still comes from the authenticated API decision. This preserves integrity while making feedback softer.
- Reload the published 0.8.20 extension ZIP in Chrome, then repeat the ten-link browser flow to verify answer-grading-v1 dispatch in the installed profile.
- For production release confidence, keep the blocked-source fixtures in the regression set and require a replacement grounded objective rather than relaxing the quality gates.

Release decision: code and deployment checks pass; question-quality QA is safer and more honest, with 8/10 sources accepted and 2/10 blocked for repair instead of shipping defects.