

ClipQuest post-fix live acceptance report

19 August 2026 - ten-link Chrome evidence, native Android live interaction, iOS baseline, AI quiz generation, answer flows, grading, question quality, and cheat-sheet export

Executive result

The current release is deployed and the AI-only boundary is intact: quiz generation is local through the updated Chrome extension or native DeepSeek configuration, while the Worker stores metadata and artifacts. The post-fix Chrome pass produced several complete five-question quizzes with reasoned grading and ready cheat sheets. A fresh Android run reached real questions and reason-first grading, but its tenth local generation call was interrupted. The iOS evidence available in this turn is a ten-link baseline from the previous native binary, not acceptance evidence for commit 2456fa0. Therefore the requested ten videos on every platform are not green yet.

Surface	Coverage	Best verified result	Current blocker	Decision
Chrome web	10 distinct YouTube IDs; updated extension 0.8.22	5 best-run banks reached 5/5; MC/TF/short answers completed; AI reasons visible; completion state and Export notes enabled	Several concurrent tabs still show network interruptions, one invalid generation event, and an automated browser download-capture gap	AMBER - repeat after cleanup
Android emulator	1 fresh signed-in QA account; IETP source exercised	q1 correct; q2 incorrect with a soft reason; 9/10 generation calls accepted	final call abandoned as network_interrupted; quiz remained on retrying and never reached completion/export	RED - P0 transport/recovery
iPhone 17 Pro simulator	10-link native baseline exists, but not rebuilt/re-run against 2456fa0	historical baseline reached imported videos and some graded questions	old binary showed partial banks, protocol mismatch, short-answer failures, and no native Library export action	BLOCKED - fresh current-build run required

Release evidence

- Source: main and origin/main at 2456fa0. The patch adds an active local-generation lease guard so fresh DeepSeek calls are not incorrectly abandoned by automatic recovery.
- Cloudflare: Worker version c5fe540f-2649-4fb2-be83-c96fdbd249d9, tag 2456fa0, promoted to 100%. Health returned ok=true, prompt quiz-local-json-stream-v5.12, validator v5.3, storage-only/local-extension generation, and request-key affinity.
- Extension: ClipQuest Local AI 0.8.22, unpacked build reloaded in Chrome; ZIP SHA-256 8e4f240d100439ec7c8f44828091f8ef5ace75369f9d5968d6b907444a0635d8.
- Automated verification: root typecheck, full API suite (27 files / 179 tests), app tests (34 files / 141 tests), extension tests (111/111), full build, and web asset verification passed before live QA.

Chrome: ten-link live matrix

The runner used production ClipQuest, all three question types, Yes, quiz me, and Short - 5. Rows report the best verified post-fix run for each link; lingering tabs also generated separate telemetry, so a single noisy tab is not presented as the product average.

#	YouTube ID	Best generation result	Answer / grading evidence	Cheat sheet
1	dEFtrYG9P40	2/5 in latest tab; an earlier run reached q2	MC correct; short answer accepted in one run with an AI reason	Not fresh-ready in latest tab
2	ZwnL06IfK_0	5/5; D1 bank 36b9bd06... passed; retryCount 0	MC, TF, and short-answer path completed with stored reasons	Ready metadata observed
3	STWIAcZfyK0	5/5	All five completed; MC, TF, and short answer correct; AI/server reasons visible	Ready; sheet 3666bfb3-7aef-4887-8861-314a9211c711
4	IETP-QEBhKw	5/5	MC, TF, short answer, TF, MC; all answered correctly with reasons	Ready; D1 calls 0-4 complete
5	NytVSSjPEgA	5/5; D1 bank fdc1679b... passed; retryCount 0	q4 short answer accepted: "Round-off error, also called rounding error." with AI reason; math rendered	Ready; sheet 2631e259-6fc0-4d11-b390-422652e44140
6	xWGMp4nyQtA	1/5 latest; an earlier run completed with retry	MC correct; later run hit network_interrupted on call 1	Not fresh-ready in latest tab
7	_TVkBfFDUwA	5/5; D1 bank e9842810... passed; retryCount 0	Mixed types answered; completion state reached	Ready metadata observed
8	KqxEtj5VvZs	0/5; creation stopped	No question; invalid generation call event surfaced	Unavailable
9	K5DwJ1HeYJc	1/5 latest; call 1 network_interrupted	q1 MC correct; generation did not finish	Unavailable
10	Wp7xBG6QIY8	3/5 latest; call 3 network_interrupted	q1/q2 correct; q3 short answer was initially rejected with a reason, retry was accepted	Unavailable in latest tab

Chrome export check

The fresh Nyt completion screen exposed an active Export notes button and a ready private cheat-sheet record. A browser download event was not captured by the automation harness, so this is UI/server readiness evidence rather than a filesystem proof for that fresh run. Separately, an existing completed Library artifact at /Users/unoxyrich/Downloads/APCSP-U1-L8-Color-Images-cheat-sheet.pdf was inspected: 2,043 bytes, one Letter page, and legible after Poppler rendering. The web export path is therefore operational, but its automated download assertion needs a stronger wait/filename check.

Android: native live interaction

A fresh account qaandroid819@example.com was created and verified for this run. The Android emulator was booted, the debug build ran with Metro, and the DeepSeek key was configured in Settings → Local AI without exposing the key in logs or this report. The source was IETP-QEBhKw and the native default session length was ten questions.

Stage	Observed result	Evidence / impact
Setup	Home, Settings, Local AI, import, and Video ready all worked	Native UI rendered; key status changed to Configured on this device
Question 1	Multiple choice selected and checked correct	Visible feedback: Nice! That's right.
Question 2	True/false selected True and checked incorrect	Soft reason-first feedback explained that 100 patterns are 0 through 99, not 1 through 100; no harsh fallback text
Background generation	Calls 0-7 completed, accepted_count 1 each	D1 telemetry: nine of ten question slots were accepted
Final call	Call 8 abandoned as network_interrupted after approximately 120 seconds	Quiz remained at q2 retrying; no completion score or export action was reachable

Android defect assessment

- The native key and question renderer are functional; the failure is in transport/recovery at the last local-generation call, not a missing AI key.
- The learner could answer while generation continued, satisfying the non-blocking design up to the point where the current question needed the missing suffix.
- The retrying state needs a bounded, visible recovery action and must never strand a learner indefinitely after the final call is abandoned.

iOS boundary

The available ten-link iOS evidence is the prior native binary, not the 2456fa0 build. That baseline showed partial generation, a protocol-version union error on KqxEtj5VvZs, repeated short-answer no-reason decisions, and no Export notes action in Library. It is included as a known-problem inventory, not as a current-release pass. A fresh iOS build/run is still required before claiming cross-platform acceptance.

Confirmed problems to fix

ID	Priority	Confirmed live problem	Fix / acceptance test
P0-1	P0	Generation suffix calls can be abandoned as <code>network_interrupted</code> on Chrome and Android; learner stays in Preparing/Retrying and cannot complete.	Trace dispatch, stream activity, completion, and abort per ordinal. Keep the active-call lease through the actual DeepSeek timeout; add one bounded reconnect and an explicit Retry generation action. Acceptance: ten fresh links complete the selected bank or show an actionable terminal error, never an unbounded spinner.
P0-2	P0	KqxEtj5VvZs produced an invalid generation call event instead of a repaired question.	Share one v5 request/event schema between extension, native clients, and API. Validate at the bridge and log a redacted schema error. Acceptance: malformed event is rejected and only the missing suffix is retried.
P0-3	P0	Short-answer grading sometimes lacks the required AI reasoned decision.	Require DeepSeek tool output <code>{is_correct, reason}</code> ; bounded retries for malformed output; preserve quiz completion but expose retry/configuration state. Acceptance: exact, paraphrased, incomplete, and fragment answers each receive a reason before correct/incorrect.
P1-1	P1	True/false grading can infer an unstated absolute claim (for example "can" treated as "always") or flip polarity after retry.	Grade the literal proposition and preserve polarity. Add qualifiers/polarity lint plus adversarial tests for can/may/only/always, negation, and zero-based ranges.
P1-2	P1	Some math prompts omit assumptions such as whether symbols may repeat; a 90-pattern item became ambiguous.	Require order/repetition assumptions in the question and validate the numeric explanation against the answer key before accepting the item.
P1-3	P1	One live short-answer prompt drifted from Black and White Images into package-tracking frequency.	Run concept-overlap/source-grounding validation against primer and quiz context; regenerate only the offending item.
P1-4	P1	Native Library did not expose Export notes in the baseline iOS UI; fresh iOS parity remains unverified.	Add ready/preparing/retry/disabled actions to native VideoCard and completion screen; test private download and expo-sharing on both native platforms.
P2-1	P2	Fresh Chrome Export notes was active and server-ready, but the browser harness did not observe a download event.	Wait for the actual download path, assert safe filename and file bytes, then render the PDF in CI/QA. Keep the server artifact check separate from the filesystem assertion.
P2-2	P2	Chrome has multiple lingering tabs producing concurrent telemetry and making latest-tab results noisy.	Close each isolated tab after capture, persist one run ledger per video, and label best-run versus latest-run data in the report.

Question-quality observations

- Positive: completed post-fix banks mixed multiple choice, true/false, and short answer; accepted short answers included concise paraphrases and complete fragments, and the grader supplied a reason before the decision.
- Positive: the Android q2 response demonstrates softer grading language ("Almost—try this concept another way") while still explaining the zero-based correction.
- Risk: duplicate concepts are tolerable only when the prompt and answer target differ. Exact repeat detection still belongs in the question-bank validator.
- Risk: incomplete banks must never receive a completion score or cheat-sheet export. The live partial runs correctly remained unexportable, but the learner-facing recovery needs to be more decisive.

Retest plan and release decision

- 1. Fix the v5 bridge envelope and network-interruption recovery first. Run one clean browser tab and one clean native session; prove every ordinal has a completed lifecycle event.
- 2. Exercise the AI grader with exact, paraphrased, incomplete, fragmentary, unsupported-absolute, negated, and true/false polarity cases. Assert reason-first output and a softer but stable decision.
- 3. Add source-grounding, assumption, and duplicate-target validators. Reject only the bad item and regenerate the missing suffix; do not restart a good prefix.
- 4. Rebuild the iOS binary from 2456fa0 or later, then run the same ten-link ledger. Verify Android with the repaired emulator path, including completion and share/export.
- 5. Rerun Chrome, iOS, and Android independently with all three question types, answer every question, capture D1 generation events, download each ready PDF, and render-check the bytes. Only then call the release green.

Final decision

Not green for the requested scope. The current deployment and AI-only architecture are healthy enough for focused repair, and Chrome has credible post-fix full-run evidence. The remaining P0 generation/transport and grading-envelope failures still prevent a reliable ten-video, three-platform learner journey. This report intentionally records the blocked platform coverage rather than inventing completions or PDF exports.